



ConvNeXtV2 Tiny Food-101 Full Training Evidence

10 epochs | batch size 16 | accumulation 1 | learning rate 0.0001 | CUDA RTX 4060

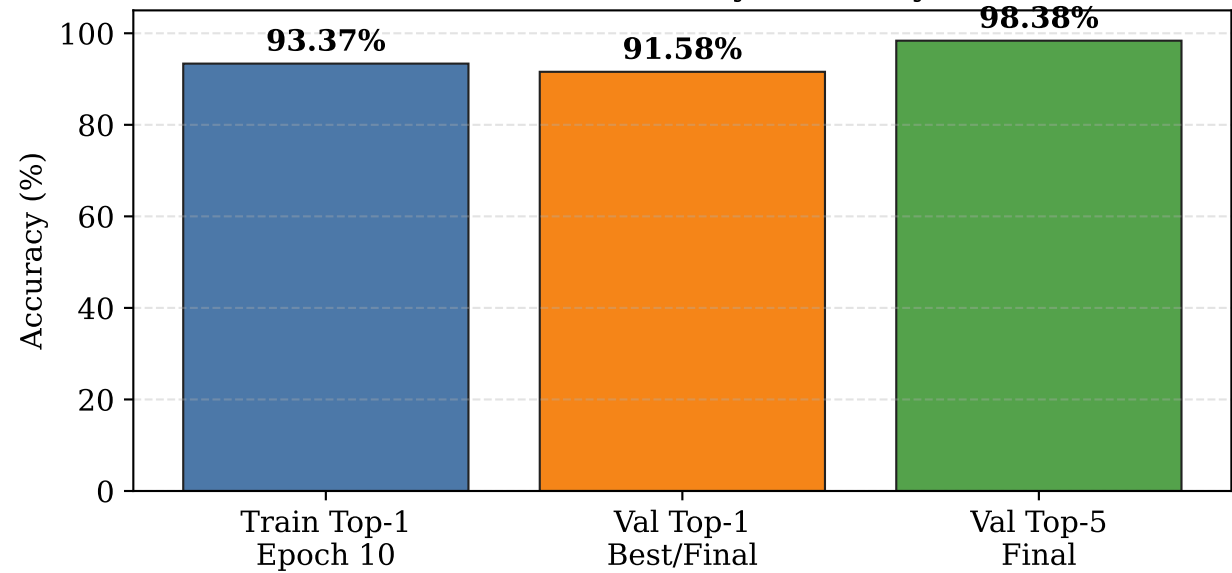
(a) Training and Validation Loss



(b) Accuracy Progression



(c) Final Accuracy Summary



(d) Benchmark Run Details

Model	ConvNeXtV2 Tiny
Dataset	Food-101
Train samples	75,750
Validation/test samples	25,250
Classes	101
Training time	525.83 minutes
Best validation top-1	91.58%
Final validation top-5	98.38%

This figure replaces the earlier smoke-run visual evidence with the completed 10-epoch benchmark run.

Figure: Visual evidence from the completed ConvNeXtV2 Tiny Food-101 training run. Validation top-1 improved from 82.64% to 91.58%, while final top-5 accuracy reached 98.38%.